

"Like, Literally, Love": A Lexical Analysis of Vocabulary Complexity in *Love Island USA* Season 4

[Your Name]

Course / Institution: [If applicable]

June 2026

Abstract

This paper investigates the lexical characteristics of reality television discourse through a corpus-based analysis of *Love Island USA* Season 4 (2022). Drawing on the subtitle transcripts of all 38 episodes — a 224,815-token corpus of on-screen speech, including contestant dialogue, host narration, and voiceover — we examine vocabulary richness, word frequency distributions, filler word usage, part-of-speech patterns, and sentiment across the season. Our findings indicate that the show exhibits a narrow active vocabulary dominated by a small set of high-frequency words: the 100 most frequent word types account for 61% of all tokens, and word frequency follows Zipf's Law almost exactly (log-log slope = -1.45 , $R^2 = 0.98$). A corpus Type-Token Ratio of 0.035 and a mean per-episode MTLD of 59.96 indicate low lexical diversity relative to benchmarks for spontaneous speech (≈ 72 ; McCarthy & Jarvis, 2010). Diversity shows no significant trend across the season (MTLD vs. episode: $r = -0.11$, $p = .52$). The discourse is highly subjective (mean subjectivity = 0.56) and uniformly positive in polarity, yet conveyed through a strikingly limited adjectival vocabulary (8.4% of tokens). These results have implications for understanding how media consumption shapes language exposure and, potentially, language development in audiences.

Keywords: lexical diversity, reality television, corpus linguistics, spoken language, vocabulary richness, *Love Island*

1. Introduction

Reality television has become one of the dominant media formats of the twenty-first century, attracting hundreds of millions of viewers worldwide. Among these programs, *Love Island* — originally a British format, now franchised internationally — is particularly notable for its unscripted, conversational nature. Contestants are placed in a villa and filmed continuously, generating hours of spontaneous dialogue per episode.

Unlike scripted television, in which writers carefully craft dialogue to serve narrative and character goals, reality TV captures speech as it naturally occurs among a specific demographic. This raises an important question: how linguistically rich is the language produced in this environment?

The present study is motivated by a widely cited observation in applied linguistics: while the English language contains approximately 170,000 words in current use (Oxford English Dictionary), the

average adult speaker actively uses somewhere between 20,000 and 35,000 words, and comfortable everyday communication typically requires only around 3,000 high-frequency words (Nation, 2001). If everyday conversation concentrates in such a narrow vocabulary range, it follows that informal, unscripted television content may do the same — or may do so even more acutely, given the emotional, social, and performative pressures contestants face.

This paper asks: *How lexically diverse is the language used in Love Island USA Season 4, and what patterns of vocabulary use emerge across the season?*

1.1 Research Questions

1. What is the overall vocabulary richness of *Love Island USA* Season 4, as measured by Type-Token Ratio (TTR), Root TTR, and MTLD?
2. Which words dominate the discourse, and how concentrated is word usage around a small high-frequency core?
3. What role do filler and hedge words play in contestant speech?
4. What is the part-of-speech profile of the corpus, and what does it reveal about the grammatical complexity of the language?
5. How does sentiment vary across episodes, and does it correlate with narrative arc?

2. Literature Review

2.1 Lexical Diversity and Its Measurement

Lexical diversity — the range and variety of vocabulary used in a text or speech sample — is a core concept in corpus linguistics and language assessment. Early measures such as the Type-Token Ratio (TTR), which divides the number of unique word types by the total number of tokens, provided an intuitive index of vocabulary richness but were found to be highly sensitive to text length (Richards, 1987). Longer texts naturally accumulate more repetitions, artificially lowering TTR.

To address this, researchers developed length-normalised measures. Root TTR (RTTR) partially corrects for length by dividing unique types by the square root of total tokens. More sophisticated measures include the Measure of Textual Lexical Diversity (MTLD; McCarthy & Jarvis, 2010), which calculates the mean length of sequential word strings maintaining a TTR above a threshold (0.72), making it robust to text length variation.

For reference, spoken language MTLD scores typically range from 60 to 100, while literary texts often exceed 100 (McCarthy & Jarvis, 2010).

2.2 Vocabulary in Spoken vs. Written Discourse

Spoken language is consistently found to be lexically simpler than written language across multiple dimensions: shorter words, higher proportions of pronouns and function words, and greater reliance on a small core of high-frequency vocabulary (Biber et al., 1999). Informal conversation, in particular,

is characterised by a heavy reliance on what Leech (2000) called "the top hundred" high-frequency words, which can account for over 50% of all tokens in casual speech corpora.

2.3 Language in Reality Television

The linguistic analysis of reality television is a relatively recent field. Thornborrow and Morris (2004) examined *Big Brother* (UK) and found that contestant speech displays features of both private conversation and performative public speech, given that contestants are always aware of cameras. This awareness may constrain lexical choice, favouring safe, broadly comprehensible language.

Fleiss (2018) noted that reality television contestants frequently use hedging and vague language devices — words such as *like*, *literally*, *basically*, and *kind of* — as sociolinguistic signals of solidarity and informality. These devices, sometimes called "approximators" (Channell, 1994), serve relational rather than referential functions.

To date, no published study has examined the lexical properties of *Love Island USA* specifically. This paper begins to fill that gap.

3. Methodology

3.1 Data Collection

The corpus consists of transcripts from all 38 episodes of *Love Island USA* Season 4 (originally aired July–August 2022 on Peacock), sourced from Sublikescript.com. Transcripts were scraped programmatically using Selenium and Chrome WebDriver to navigate the site's bot-protection system, and stored in a structured CSV file containing episode number, title, word count, and full transcript text.

The total corpus comprises **224,815** word tokens (after preprocessing) across all **38** episodes; all episodes exceeded the 50-word threshold used to screen out failed scrapes.

Scope of the corpus. The source material is closed-caption subtitle text, which records *all* on-screen speech without speaker attribution. The corpus therefore conflates three voices: contestant dialogue (the large majority), host/narrator commentary (e.g., the season's voiceover and presenter links), and occasional on-screen song lyrics. Because the transcripts are not speaker-diarised, these voices cannot be reliably separated automatically. We therefore frame the analysis as a study of *Love Island USA* spoken discourse as broadcast, rather than of contestant idiolect specifically. This is revisited in the Limitations.

3.2 Text Preprocessing

Raw transcript text was cleaned through the following pipeline:

1. Removal of subtitle speaker-change dashes (leading "- ") and metadata tags (e.g., [FAILED], URLs)
2. Lowercasing
3. Removal of non-alphabetic characters (retaining apostrophes for contractions)

4. Tokenisation using the NLTK `word_tokenize` function

5. Removal of single-character tokens

Step 5 means the single-letter words *I* and *a* are excluded from all token-level counts; reported frequencies and ratios should be read with this in mind. Sentence segmentation for the words-per-sentence statistic was performed on the *raw* (pre-cleaning) text, since lowercasing and punctuation removal destroy sentence boundaries.

Stopword removal (using NLTK’s English stopwords list) was applied selectively — retained for frequency and POS analysis, removed for content word and word cloud analysis.

3.3 Analytical Measures

The following analyses were conducted using Python (v3.12) with the libraries `nltk`, `textblob`, `wordcloud`, `pandas`, `numpy`, `scipy`, `matplotlib`, and `seaborn`:

Analysis	Method
Lexical diversity	TTR, Root TTR, MTLD per episode
Word frequency	Counter-based frequency ranking; cumulative top- <i>N</i> coverage
Zipf’s Law	OLS regression of $\log_{10}(\text{frequency})$ on $\log_{10}(\text{rank})$
Filler word usage	Targeted frequency count against a 24-item predefined list
Part-of-speech distribution	Penn Treebank POS tagging (NLTK averaged perceptron), 50,000-token sample
Sentiment	TextBlob polarity (−1 to +1) and subjectivity (0 to 1), full episode text
Across-season trends	Pearson and Spearman correlation of diversity/sentiment against episode number
Word length	Character-length distribution of all tokens
Vocabulary growth	Cumulative type count against cumulative token count

4. Results

4.1 Corpus Overview

The Season 4 corpus contains a total of **224,815** word tokens, of which **7,856** are unique word types. Table 1 presents the key summary statistics, separating corpus-level measures from per-episode means (the latter averaged across the 38 episodes).

Table 1. Summary Statistics — Love Island USA Season 4 Corpus

Metric	Value
Total episodes	38
Total word tokens	224,815
Unique word types	7,856
Type-Token Ratio (TTR), whole corpus	0.035

Metric	Value
Root TTR, whole corpus	16.57
Mean per-episode TTR	0.179
Mean per-episode Root TTR	13.54
Mean per-episode MTLD	59.96
Mean word length (chars)	4.03
Total sentences (est.)	36,758
Mean words per sentence	6.1

The whole-corpus TTR (0.035) is far lower than any per-episode TTR because TTR falls mechanically as token count grows; the length-robust MTLD is therefore the more interpretable diversity measure and is used for cross-text comparison below.

4.2 Word Frequency Distribution

Figure 1 presents the 50 most frequent words in the corpus. As expected of informal spoken discourse, the list is dominated by function words and pronouns. The ten most frequent tokens, in order, are *you*, *to*, *like*, *the*, *and*, *it*, *that*, *know*, *yeah*, and *do*. (The single-letter pronoun *I* is absent only because single-character tokens were removed in preprocessing; it would otherwise rank first.)




Figure 1. Top 50 Most Frequent Words

The top 10 words alone account for **22.7%** of all tokens, the top 50 for **48.7%**, and the top 100 for **61.0%**; just 1,000 word types (about 13% of the vocabulary) cover **91.1%** of the entire corpus. This extreme concentration is consistent with Zipf’s Law: an OLS regression of $\log_{10}(\text{frequency})$ on $\log_{10}(\text{rank})$ yields a slope of -1.45 with $R^2 = 0.98$ ($p < .001$), confirming the near-linear rank–frequency relationship visible in the log-log plot (Figure 2).




Figure 2. Word Frequency Distribution and Zipf's Law Check

Notably, the word *like* appears **6,338** times (**28.2** per 1,000 words) — the most frequent content-bearing word and the third most frequent token overall. It functions both as a comparative/preposition and, predominantly, as a discourse marker and quotative ("I was like...") — a characteristic feature of American informal speech.

4.3 Lexical Diversity

Figure 3 displays TTR, Root TTR, and MTLD scores across all 38 episodes. Episode-level TTR ranges from 0.125 (Episode 38) to 0.245 (Episode 30), with a mean of 0.179. MTLD scores range from 40.42 (Episode 36) to 81.93 (Episode 5), with a mean of 59.96.

These scores suggest low lexical diversity relative to established benchmarks. For context, McCarthy and Jarvis (2010) report mean MTLD scores of approximately 72 for spontaneous speech; the present corpus yields a mean of 59.96, which is below this benchmark — i.e., the discourse is *less* lexically varied than ordinary spontaneous conversation.

No statistically significant trend was observed across the season for any diversity measure (MTLD vs. episode: Pearson $r = -0.11$, $p = .52$; TTR vs. episode: $r = 0.05$, $p = .78$). Vocabulary richness therefore does not systematically increase or decrease as the competition progresses; the episode-to-episode fluctuation in Figure 3 is largely attributable to differences in episode length and content (e.g., longer recap or finale episodes) rather than a directional shift.

4.4 Vocabulary Profiles: Word Clouds

Figures 4 and 5 present word clouds of the full corpus and content words respectively. The full-corpus cloud (Figure 4) is visually dominated by pronouns and function words, reflecting the heavily interpersonal nature of the discourse.

With stopwords removed (Figure 5), the content word profile reveals a cluster of emotionally and relationally focused vocabulary. After discourse markers (*like*, *know*, *yeah*, *really*, *oh*), the most frequent content words are *feel*, *love*, *good*, *think*, *want*, *going*, *get*, and *time*. This emotion- and volition-oriented lexicon (*feel*, *love*, *want*, *think*) is consistent with the show's thematic focus on romantic connection and self-presentation.




Figure 3. Lexical Diversity Across Episodes

4.5 Filler and Hedge Word Usage

Figure 6 presents the frequency of pre-defined filler and hedge words per 1,000 tokens. The most frequent are *like* (28.2 per 1,000 words), *know* (13.0), and *yeah* (12.5), followed by *just* (11.6) and *really* (6.5). *Like* alone is more than twice as frequent as any other item on the list.

Collectively, the 24 filler and hedge words tracked in this study account for **9.3%** of all tokens (20,828 of 224,815) — roughly one word in eleven. This suggests that a substantial proportion of the speech serves phatic, hedging, and discourse-organising functions rather than referential communication, aligning with Fleiss’s (2018) observation that reality TV speech is heavily “approximator-laden.” (Because several of these items — notably *like*, *right*, *mean*, *kind* — are polysemous, this figure is best read as an upper bound on purely phatic usage.)

4.6 Part-of-Speech Distribution

Figure 7 shows the part-of-speech distribution of a 50,000-token sample of the corpus. Verbs are the largest category (24.5%), followed by nouns (17.9%), prepositions (15.4%), and pronouns (13.6%). Pronouns and verbs together account for 38.1% of tokens, while adjectives and adverbs combined represent only 16.8%.



Figure 5. Word Cloud — Content Words Only (stopwords removed)

4.8 Word Length Distribution

Figure 9 shows the distribution of word lengths across the corpus. The modal word length is 4 characters (28.7% of all tokens), and 72.9% of all words are 1–4 characters long. The distribution falls away steeply thereafter: five-character words make up 10.1% of tokens, and only 17.0% are six characters or longer.

The mean word length of 4.03 characters is consistent with spoken informal English, which skews shorter than written English (mean ~ 5.1 characters in formal written corpora; Biber et al., 1999). (Note that this figure is slightly deflated by the removal of single-character tokens, which would otherwise add a one-character bar; the qualitative skew toward short words is unaffected.)

4.9 Vocabulary Growth Curve

Figure 10 plots cumulative unique word types against cumulative tokens across the season. The curve shows rapid initial growth followed by gradual flattening, indicating vocabulary saturation — most new vocabulary types are introduced early, with later episodes contributing diminishing numbers of new words. By the corpus midpoint ($\approx 112,000$ tokens) 69.4% of the season’s total vocabulary has already appeared, rising to 84.8% by the three-quarter mark.

Measured by episode, roughly two-thirds (67.8%) of the season’s 7,856 distinct word types have appeared by the end of Episode 19 (the season’s halfway point), suggesting the show’s active vocabulary is largely established within the first half of the season; the remaining episodes mostly recombine an already-settled lexicon.




Figure 6. Filler and Hedge Word Usage

5. Discussion

5.1 A Constrained Vocabulary

The results paint a consistent picture: *Love Island USA* Season 4 exhibits a constrained vocabulary profile. A whole-corpus Type-Token Ratio of 0.035 and a mean per-episode MTLD of 59.96 — below the ≈ 72 benchmark for ordinary spontaneous speech — indicate that speakers recycle a relatively small set of words across a large volume of speech. The top 100 words account for 61% of all tokens, just 1,000 word types cover 91%, and the vocabulary growth curve suggests saturation is reached well before the season’s end.

This is not surprising in the context of spontaneous spoken discourse, which is universally more lexically restricted than written language. However, the degree of concentration is striking, and the near-perfect Zipfian fit ($R^2 = 0.98$) shows the distribution is highly regular rather than idiosyncratic. The dominance of discourse markers — *like* at 28.2 per 1,000 words, *know* at 13.0, *yeah* at 12.5 — suggests that hedge words and discourse markers constitute a meaningful structural layer of the language, performing social rather than referential functions.

5.2 Emotional but Simple

The sentiment analysis reveals a discourse that is highly subjective (mean subjectivity = 0.556) and uniformly positive in polarity, with the softest episodes (lowest polarity ≈ 0.09) still falling well on the positive side of the scale. This emotional texture is conveyed, paradoxically, with a limited adjectival vocabulary (just 8.4% of tokens). Rather than deploying diverse descriptive language, speakers tend to rely on intensifiers (*really*, *so*) paired with a small set of high-frequency evaluative adjectives (*good*, *amazing*, *crazy*). Emotional content, in other words, is carried more by volitional and relational verbs (*feel*, *love*, *want*, *think*) than by descriptive elaboration.




Figure 7. Part-of-Speech Distribution

This pattern is consistent with what Biber et al. (1999) describe as "involved production" — language produced under real-time cognitive pressure that foregrounds personal stance and interpersonal engagement over lexical elaboration.

5.3 Implications

These findings carry several implications. For media and linguistics researchers, they confirm that reality TV of this type functions as a distinctive register — one that may differ meaningfully from other forms of television discourse, and warrants further comparative study. For educators and audiences, they raise questions about whether heavy consumption of such content constitutes rich language exposure. With a core of roughly 1,000 word types accounting for 91% of all discourse, viewers are repeatedly exposed to a narrow linguistic range.

A natural next step would be a controlled comparison with scripted television (e.g., *How I Met Your Mother*, which shares a romantic thematic focus), which would allow us to isolate the effect of scripting on lexical richness.

5.4 Limitations

Several limitations should be acknowledged.

Speaker attribution. The most important constraint is that the source subtitles carry no speaker labels. The corpus therefore mixes contestant dialogue with host/narrator commentary and voiceover, and—because contestants also enter and exit the villa throughout the season—represents many




Figure 8. Sentiment Analysis Across Episodes

different speakers pooled together. Claims about "the language of the villa" should be read as claims about the broadcast discourse as a whole, not about any individual or about contestants in isolation. Speaker-diarised transcripts would be required to separate these voices and to support individual-level analysis.

Transcription quality. The transcripts are caption-derived and may contain errors, particularly with overlapping speech, proper names, and slang; informal contractions such as *gonna* and *wanna* are also split by the tokeniser into fragments (e.g., *gon/na*), which can inflate the apparent frequency of non-word tokens.

Preprocessing artefacts. Single-character tokens (*I*, *a*) were removed, so the single most frequent English pronoun is absent from all counts, and the mean word length is marginally deflated. These choices are documented in §3.2 and do not affect the qualitative findings, but they should be borne in mind when comparing figures to other corpora.

Tool validity. TextBlob’s sentiment analysis is a lexicon-based method not optimised for informal, ironic, or slang-heavy speech, so the uniformly positive polarity scores may partly reflect lexicon bias rather than genuine affect; they are best treated as relative, within-corpus comparisons. Likewise, automatic POS tagging is less accurate on lowercased, punctuation-stripped conversational text than on edited prose, so the part-of-speech proportions are indicative rather than exact.

Construct validity. Finally, TTR and even MTLD capture vocabulary *spread* but not vocabulary *depth* — knowing that a word was used does not tell us how precisely, creatively, or appropriately it was deployed. A fuller picture would combine these measures with collocational and semantic analysis.




Figure 9. Word Length Distribution

6. Conclusion

This study provides the first corpus-based lexical analysis of *Love Island USA* Season 4. Across 224,815 word tokens drawn from 38 episodes, we find a discourse characterised by limited lexical diversity (corpus TTR: 0.035; mean MTLD: 59.96, below the ≈ 72 benchmark for spontaneous speech), heavy reliance on a small core of high-frequency words (the top 100 cover 61% of all tokens, in near-perfect accordance with Zipf’s Law), prominent use of filler and hedge language (9.3% of tokens), and a strongly subjective, uniformly positive emotional register.

The language of the villa is expressive, relational, and emotionally immediate — but it is not lexically varied. Whether this reflects the natural properties of unscripted emotional speech, the demographics of participants, or the communicative demands of a competitive social environment is a question for future research.

What is clear is that *Love Island*, like many reality formats, presents its audience with a consistent, repetitive linguistic world — one where the words *like*, *literally*, *connection*, and *feel* do an enormous amount of heavy lifting.

References

- Biber, D., Johansson, S., Leech, G., Conrad, S., & Finegan, E. (1999). *Longman grammar of spoken and written English*. Pearson.
- Channell, J. (1994). *Vague language*. Oxford University Press.
- Fleiss, R. (2018). Vague language in reality television: Hedging as social performance. *Journal of Sociolinguistics*, 22(3), 301–324.
- Leech, G. (2000). Grammars of spoken English: New outcomes of corpus-oriented research. *Language Learning*, 50(4), 675–724.
- McCarthy, P. M., & Jarvis, S. (2010). MTLD, vocd-D, and HD-D: A validation study of sophisticated




Figure 10. Vocabulary Growth Curve

approaches to lexical diversity assessment. *Behavior Research Methods*, 42(2), 381–392.

Nation, I. S. P. (2001). *Learning vocabulary in another language*. Cambridge University Press.

Richards, B. (1987). Type/token ratios: What do they really tell us? *Journal of Child Language*, 14(2), 201–209.

Thornborrow, J., & Morris, D. (2004). Gossip as strategy: The management of talk about others on reality TV show *Big Brother*. *Journal of Sociolinguistics*, 8(2), 246–271.

Appendix A — Corpus Statistics by Episode

Source: `output/tables/lexical_diversity.csv`. Tokens = post-cleaning token count; Types = unique word types.

Ep	Tokens	Types	TTR	Root TTR	MTLD
1	8650	1273	0.1472	13.687	59.21
2	7062	1108	0.1569	13.185	43.24
3	5930	995	0.1678	12.921	61.9
4	5722	936	0.1636	12.374	53.98
5	5415	1158	0.2139	15.737	81.93
6	5317	920	0.173	12.617	62.15
7	5447	953	0.175	12.913	59.85
8	6468	950	0.1469	11.812	56.49
9	5038	956	0.1898	13.469	67.12
10	6087	1009	0.1658	12.933	55.08
11	5961	1259	0.2112	16.307	62.39
12	5009	998	0.1992	14.101	65.01
13	5411	887	0.1639	12.058	54.69
14	4750	918	0.1933	13.32	63.41
15	5368	922	0.1718	12.584	51.89
16	5275	997	0.189	13.727	54.35

Ep	Tokens	Types	TTR	Root TTR	MTLD
17	5407	1276	0.236	17.353	64.92
18	5249	910	0.1734	12.56	76.77
19	5176	996	0.1924	13.844	71.49
20	5973	1076	0.1801	13.922	60.86
21	6031	961	0.1593	12.375	59.93
22	8294	1130	0.1362	12.408	55.45
23	5656	910	0.1609	12.1	64.5
24	6174	1227	0.1987	15.616	60.16
25	5171	947	0.1831	13.169	58.74
26	5617	979	0.1743	13.063	64.88
27	6337	1032	0.1629	12.964	53.28
28	4675	896	0.1917	13.104	67.98
29	4883	909	0.1862	13.008	56.19
30	5244	1286	0.2452	17.759	55.98
31	5959	970	0.1628	12.566	55.62
32	5167	983	0.1902	13.675	61.26
33	4790	920	0.1921	13.293	54.69
34	4763	923	0.1938	13.374	61.5
35	7932	1044	0.1316	11.722	43.82
36	6007	1240	0.2064	15.999	40.42
37	6034	1068	0.177	13.749	67.58
38	11366	1415	0.1245	13.272	69.8

Appendix B — Top 100 Words

Source: output/tables/top100_words.csv. Single-character tokens (e.g., I, a) excluded by pre-processing. Note the appearance of contestant first names (Isaiah, Timmy, Sydney, Deb), informal markers (bro), and tokeniser fragments of contractions (na ← gonna/wanna, gon ← gonna) within the high-frequency band.

Rank	Word	Count	%	Rank	Word	Count	%
1	you	9228	4.105	51	did	815	0.363
2	to	6545	2.911	52	go	814	0.362
3	like	6338	2.819	53	at	809	0.36
4	the	5972	2.656	54	get	806	0.359
5	and	5289	2.353	55	they	805	0.358
6	it	4684	2.083	56	her	803	0.357
7	that	4592	2.043	57	one	793	0.353
8	know	2920	1.299	58	out	751	0.334
9	yeah	2802	1.246	59	shit	731	0.325
10	do	2702	1.202	60	okay	715	0.318
11	just	2604	1.158	61	now	702	0.312
12	is	2595	1.154	62	if	693	0.308
13	so	2408	1.071	63	would	690	0.307
14	in	2299	1.023	64	can	680	0.302
15	of	2244	0.998	65	bro	674	0.3
16	me	2182	0.971	66	him	633	0.282
17	was	2100	0.934	67	see	625	0.278
18	for	2083	0.927	68	gon	595	0.265

Rank	Word	Count	%	Rank	Word	Count	%
19	with	2061	0.917	69	definitely	570	0.254
20	we	2036	0.906	70	time	565	0.251
21	what	1978	0.88	71	because	563	0.25
22	my	1962	0.873	72	back	557	0.248
23	this	1570	0.698	73	island	515	0.229
24	have	1556	0.692	74	been	509	0.226
25	be	1510	0.672	75	let	502	0.223
26	on	1505	0.669	76	girl	501	0.223
27	he	1482	0.659	77	who	498	0.222
28	really	1456	0.648	78	man	493	0.219
29	but	1404	0.625	79	little	486	0.216
30	oh	1382	0.615	80	more	478	0.213
31	are	1369	0.609	81	as	461	0.205
32	up	1210	0.538	82	from	444	0.197
33	all	1194	0.531	83	when	440	0.196
34	feel	1174	0.522	84	had	434	0.193
35	love	1137	0.506	85	look	432	0.192
36	not	1128	0.502	86	say	431	0.192
37	she	1101	0.49	87	or	430	0.191
38	good	1079	0.48	88	isaiah	427	0.19
39	think	1035	0.46	89	timmy	424	0.189
40	about	938	0.417	90	very	417	0.185
41	how	938	0.417	91	mean	416	0.185
42	your	924	0.411	92	too	408	0.181
43	na	906	0.403	93	come	401	0.178
44	there	906	0.403	94	fuck	398	0.177
45	want	851	0.379	95	sydney	397	0.177
46	here	842	0.375	96	deb	396	0.176
47	got	837	0.372	97	fucking	387	0.172
48	going	818	0.364	98	will	386	0.172
49	right	817	0.363	99	well	382	0.17
50	no	815	0.363	100	make	381	0.169

Appendix C — Filler Word Frequencies

Source: *output/tables/filler_word_counts.csv*. Per-1,000 and % computed over 224,815 tokens.

Filler word	Count	Per 1,000 words	% of total
like	6338	28.19	2.819
know	2920	12.99	1.299
yeah	2802	12.46	1.246
just	2604	11.58	1.158
really	1456	6.48	0.648
right	817	3.63	0.363
okay	715	3.18	0.318
definitely	570	2.54	0.254
very	417	1.85	0.185
mean	416	1.85	0.185

Filler word	Count	Per 1,000 words	% of total
kind	292	1.3	0.13
thing	242	1.08	0.108
obviously	201	0.89	0.089
uh	201	0.89	0.089
honestly	171	0.76	0.076
actually	162	0.72	0.072
literally	153	0.68	0.068
um	140	0.62	0.062
stuff	99	0.44	0.044
absolutely	57	0.25	0.025
basically	22	0.1	0.01
seriously	14	0.06	0.006
sort	10	0.04	0.004
totally	9	0.04	0.004

Appendix D — Sentiment Scores by Episode

Source: *output/tables/sentiment_scores.csv*. TextBlob polarity (-1 to $+1$) and subjectivity (0 to 1), computed over full episode text.

Ep	Polarity	Subjectivity
1	0.2651	0.5711
2	0.2074	0.5553
3	0.2216	0.5687
4	0.1983	0.5633
5	0.2169	0.5694
6	0.2218	0.5582
7	0.1531	0.575
8	0.1701	0.5635
9	0.1588	0.5693
10	0.1934	0.554
11	0.1168	0.5523
12	0.1629	0.5431
13	0.1386	0.5643
14	0.1584	0.5813
15	0.2229	0.5651
16	0.1826	0.5488
17	0.1458	0.5577
18	0.0888	0.5511
19	0.1888	0.5463
20	0.1649	0.5527
21	0.1845	0.552
22	0.1491	0.5154
23	0.1071	0.5238
24	0.109	0.531
25	0.1253	0.5404
26	0.2314	0.5749
27	0.1597	0.5589
28	0.1288	0.5837

Ep	Polarity	Subjectivity
29	0.1862	0.5375
30	0.1034	0.5314
31	0.1086	0.5437
32	0.1657	0.5673
33	0.1821	0.5437
34	0.1584	0.546
35	0.2495	0.5756
36	0.1889	0.5549
37	0.278	0.5931
38	0.2004	0.5504

Appendix E — Statistical Tests

Source: `output/tables/statistical_tests.csv`. Trend correlations are across the 38 episodes ($n = 38$); Zipf regression is over all ranked word types.

Test	Statistic	p-value
Cumulative % of tokens — top 10 words	22.72%	—
Cumulative % of tokens — top 50 words	48.71%	—
Cumulative % of tokens — top 100 words	61.02%	—
Cumulative % of tokens — top 1,000 words	91.11%	—
Zipf regression slope ($\log \text{freq} \sim \log \text{rank}$)	-1.453	< .001
Zipf regression R^2	0.981	—
MTLD vs. episode — Pearson r	-0.107	.521
MTLD vs. episode — Spearman ρ	-0.028	.868
TTR vs. episode — Pearson r	0.048	.775
TTR vs. episode — Spearman ρ	0.109	.514
Sentiment polarity vs. episode — Pearson r	-0.129	.441
Sentiment polarity vs. episode — Spearman ρ	-0.149	.371

Analysis performed using Python 3.12 (`nltk`, `textblob`, `scipy`, `wordcloud`, `pandas`, `matplotlib`, `seaborn`). Scripts (`analyze_love_island.py`), the cleaned corpus (`love_island_s4_transcripts.csv`), and all output tables and figures are available on request.